




# Distributed Data Processing

---

- Introduction
- Distributed DBMS Architecture
- Distributed DB Design
- Semantic Data Control
- **Query Processing**
- Transaction Management


# 6. Query Decomposition and Data Localization

## ■ Query Decomposition

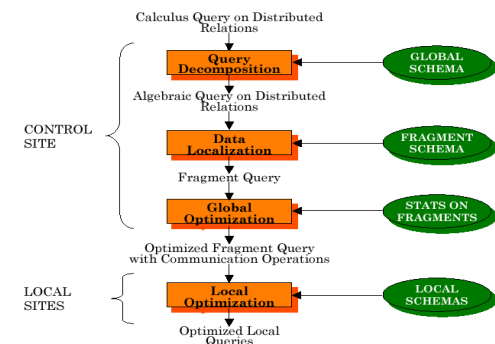


- Transform a relational *calculus query* into an *algebra query* on global relation

## ■ Localization of Distributed Data



- *Localize* the query's data using data distributed information





---

## 6.1 Query Decomposition




# Query Decomposition

---

**Input** : relational **calculus query** on global relations

- Normalization
  - manipulate query quantifiers and qualification
- Analysis
  - detect and reject “**incorrect**” queries
  - possible for only a subset of relational calculus
- Simplification
  - eliminate **redundant** predicates
- Restructuring
  - calculus query  $\Rightarrow$  algebraic query
  - more than one translation is possible
  - use transformation rules

**Output**: relational **algebra query** on global relations





## 6.1.1 Normalization

---

- Transform the query to a **normalized form** to facilitate further processing

- Put into normal form

- Conjunctive normal form

$$(p_{11} \vee p_{12} \vee \dots \vee p_{1n}) \wedge \dots \wedge (p_{m1} \vee p_{m2} \vee \dots \vee p_{mn})$$

- Disjunctive normal form

$$(p_{11} \wedge p_{12} \wedge \dots \wedge p_{1n}) \vee \dots \vee (p_{m1} \wedge p_{m2} \wedge \dots \wedge p_{mn})$$

OR's mapped into union

AND's mapped into join or selection



# Equivalent rules for logical operations

---

1.  $p1 \wedge p2 \Leftrightarrow p2 \wedge p1$
2.  $p1 \vee p2 \Leftrightarrow p2 \vee p1$
3.  $p1 \wedge (p2 \wedge p3) \Leftrightarrow (p1 \wedge p2) \wedge p3$
4.  $p1 \vee (p2 \vee p3) \Leftrightarrow (p1 \vee p2) \vee p3$
5.  $p1 \wedge (p2 \vee p3) \Leftrightarrow (p1 \wedge p2) \vee (p1 \wedge p3)$
6.  $p1 \vee (p2 \wedge p3) \Leftrightarrow (p1 \vee p2) \wedge (p1 \vee p3)$
7.  $\text{NOT}(p1 \wedge p2) \Leftrightarrow \text{NOT}(p1) \vee \text{NOT}(p2)$
8.  $\text{NOT}(p1 \vee p2) \Leftrightarrow \text{NOT}(p1) \wedge \text{NOT}(p2)$
9.  $\text{NOT}(\text{NOT}(p)) \Leftrightarrow p$



# Example of normalization

---

- The **conjunctive normal form** is more practical since query qualifications typically include more AND than OR predicates
- **Example:** Find the names of employees who have been working on project P1 for 12 or 24 months


```
SELECT  ENAME  
FROM    EMP, ASG  
WHERE   EMP.ENO = ASG.ENO  
AND     ASG.PNO = "P1"  
AND     DUR=12 OR DUR=24
```

**In conjunctive normal form:**

**$EMP.ENO=ASG.ENO \wedge ASG.PNO="P1" \wedge (DUR=12 \vee DUR=24)$**

**In disjunctive normal form:**

**$(EMP.ENO=ASG.ENO \wedge ASG.PNO="P1" \wedge DUR=12) \vee$   
 $(EMP.ENO=ASG.ENO \wedge ASG.PNO="P1" \wedge DUR=24)$**





## 6.1.2 Analysis

---

- Example

```
SELECT E#  
FROM EMP  
WHERE ENAME>200
```

! Undefined  
attribute

! Type  
mismatch

- Rejection of incorrect queries
- Type incorrect
  - If any of its attribute or relation names are not defined in the global schema
  - If operations are applied to attributes of the wrong type
  - Similar to type checking
- Semantically incorrect
  - Only a subset of relational calculus queries can be tested for correctness
    - Those that do not contain disjunction and negation
  - Components do not contribute in any way to the generation of the result
  - To detect
    - query graph (connection graph)
    - join graph




# Query graph

```
SELECT      ENAME, RESP
FROM        EMP, ASG, PROJ
WHERE       EMP.ENO = ASG.ENO
AND         ASG.PNO = PROJ.PNO
AND         PNAME = "CAD/CAM"
AND         DUR ≥ 36
AND         TITLE = "Programmer"
```

Query graph



Join graph



# Example of Analysis

- If the query graph is not connected, the query is wrong.
- Example

*SELECT* ENAME,RESP  
*FROM* EMP, ASG, PROJ  
*WHERE* EMP.ENO = ASG.ENO  
*AND* PNAME = "CAD/CAM"  
*AND* DUR ≥ 36  
*AND* TITLE = "Programmer"






## Solution to the incorrect problem

---

- **Reject** the query
- **Assume** that there is an implicit Cartesian product between the two relations
- **Infer** (using the schema) the missing join predicate





## 6.1.3 Simplification

---

### ■ Elimination of Redundancy

- A user query typically expressed on a view may be enriched with several predicates to **achieve view-relation correspondence**, and ensure **semantic integrity and security** – that may then contain redundant predicates

### ■ Idempotency rules

1.  $p \wedge p \Leftrightarrow p$
2.  $p \vee p \Leftrightarrow p$
3.  $p \wedge \text{true} \Leftrightarrow p$
4.  $P \vee \text{false} \Leftrightarrow p$
5.  $p \wedge \text{false} \Leftrightarrow \text{false}$
6.  $P \vee \text{true} \Leftrightarrow \text{true}$
7.  $p \wedge \text{NOT}(p) \Leftrightarrow \text{false}$
8.  $p \vee \text{NOT}(p) \Leftrightarrow \text{true}$
9.  $p1 \wedge (p1 \vee p2) \Leftrightarrow p1$
10.  $p1 \vee (p1 \wedge p2) \Leftrightarrow p1$





# Example1: Semantic Data Control – View Management

---

Query expressed on view

```
SELECT ENAME, PNO, RESP  
FROM SYSN, ASG  
WHERE SYSN.ENO = ASG.ENO
```

Query expressed on base relation

```
SELECT ENAME,PNO,RESP  
FROM EMP, ASG  
WHERE EMP.ENO = ASG.ENO  
AND TITLE = "Syst. Anal."
```

Definition of view

```
CREATE VIEW SYSN(ENO,ENAME)  
AS      SELECT ENO,ENAME  
        FROM EMP  
        WHERE TITLE="Syst. Anal."
```



## Example2: Semantic Data Control – Semantic Integrity Control

---

- Increasing the budget of CAD/CAM project by 10%

```
UPDATE PROJ  
SET BUDGET = BUDGET*1.1  
WHERE PNAME = "CAD/CAM"
```

```
UPDATE PROJ  
SET BUDGET = BUDGET*1.1  
WHERE PNAME = "CAD/CAM"  
AND NEW.BUDGET ≥ 500000  
AND NEW.BUDGET ≤ 1000000
```



$$p \wedge p \Leftrightarrow p$$

$$p \vee p \Leftrightarrow p$$

$$p \wedge \text{true} \Leftrightarrow p$$

$$p \vee \text{false} \Leftrightarrow p$$

$$p \wedge \text{false} \Leftrightarrow \text{false}$$

$$p \vee \text{true} \Leftrightarrow \text{true}$$

$$p \wedge \text{NOT}(p) \Leftrightarrow \text{false}$$

$$p \vee \text{NOT}(p) \Leftrightarrow \text{true}$$

$$p1 \wedge (p1 \vee p2) \Leftrightarrow p1$$


$$p1 \vee (p1 \wedge p2) \Leftrightarrow p1$$

## Example of simplification

<i>SELECT</i>	<i>TITLE</i>
<i>FROM</i>	<i>EMP</i>
<i>WHERE</i>	<i>EMP.ENAME</i> = "J. Doe"
<i>OR</i>	( <i>NOT</i> ( <i>EMP.TITLE</i> = "Programmer"))
<i>AND</i>	( <i>EMP.TITLE</i> = "Programmer"
<i>OR</i>	<i>EMP.TITLE</i> = "Elect. Eng.")
<i>AND</i>	<i>NOT</i> ( <i>EMP.TITLE</i> = "Elect. Eng."))



<i>SELECT</i>	<i>TITLE</i>
<i>FROM</i>	<i>EMP</i>
<i>WHERE</i>	<i>EMP.ENAME</i> = "J. Doe"





## 6.1.4 Restructuring

---

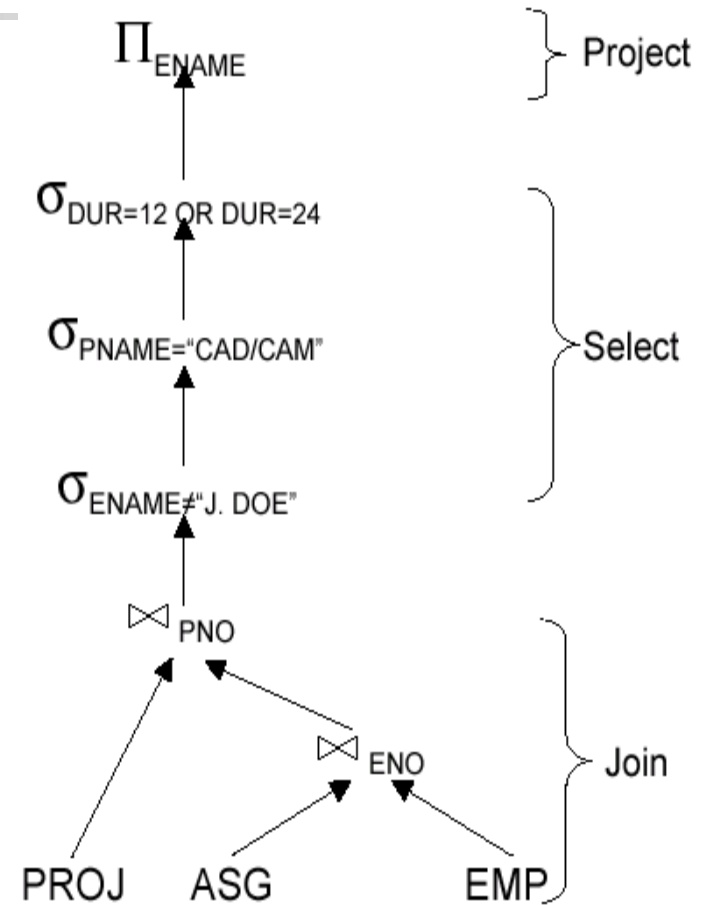
- **Rewrite** the query based on relational calculus to that based on relational algebra
  - Straightforward transformation of the query from relational calculus into relational algebra
  - Restructuring of the relational algebra query to improve performance by transformation rules
- **operator tree**



# Example of operation tree

Find the names of employees other than J. Doe who worked on the CAD/CAM project for either 1 or 2 years.

**SELECT** ENAME  
**FROM** EMP, ASG, PROJ  
**WHERE** EMP.ENO = ASG.ENO  
AND ASG.PNO = PROJ.PNO  
AND ENAME ≠ "J. Doe"  
AND PNAME = "CAD/CAM"  
AND (DUR = 12 OR DUR = 24)





# Transformation rules

---

- Commutativity of binary operations
  - $R \bowtie S \Leftrightarrow S \bowtie R$
  - $R \infty S \Leftrightarrow S \infty R$
  - $R \cup S \Leftrightarrow S \cup R$
- Associativity of binary operations
  - $(R \bowtie S) \bowtie T \Leftrightarrow R \bowtie (S \bowtie T)$
  - $(R \infty S) \infty T \Leftrightarrow R \infty (S \infty T)$
- Idempotence of unary operations
  - $\Pi_{A'}(\Pi_{A''}(R)) \Leftrightarrow \Pi_{A'}(R)$
  - $\delta_{p1(A1)}(\delta_{p2(A2)}(R)) = \delta_{p1(A1) \wedge p2(A2)}(R)$   
where  $R[A]$  and  $A'$  belongs to  $A$ ,  $A''$  belongs to  $A$  and  $A'$  belongs to  $A''$
- Commuting selection with projection
  - $\Pi_{A1 \dots An}(\delta_{p(Ap)}(R)) \Leftrightarrow \Pi_{A1 \dots An}(\delta_{p(Ap)}(\Pi_{A1 \dots An \ Ap}(R)))$



# Transformation rules

---

- Commuting selection with binary operations

- $\delta_{p(A)}(R \times S) \Leftrightarrow (\delta_{p(A)}(R)) \times S$
- $\delta_{p(A_i)}(R \bowtie_{(A_j, B_k)} S) \Leftrightarrow (\delta_{p(A_i)}(R)) \bowtie_{(A_j, B_k)} S$
- $\delta_{p(A_i)}(R \cup T) \Leftrightarrow (\delta_{p(A_i)}(R)) \cup (\delta_{p(A_i)}(T))$

where  $A_i$  belongs to  $R$  and  $T$

- Commuting projection with binary operations

- $\Pi_C(R \times S) \Leftrightarrow \Pi_{A'}(R) \times \Pi_{B'}(S)$
- $\Pi_C(R \bowtie_{(A_j, B_k)} S) \Leftrightarrow \Pi_{A'}(R) \bowtie_{(A_j, B_k)} \Pi_{B'}(S)$
- $\Pi_C(R \cup S) \Leftrightarrow \Pi_C(R) \cup \Pi_C(S)$

where  $R[A]$  and  $S[B]$ ;  $C = A' \cup B'$  where  $A'$  belongs to  $A$ ,  $B'$  belongs to  $B$






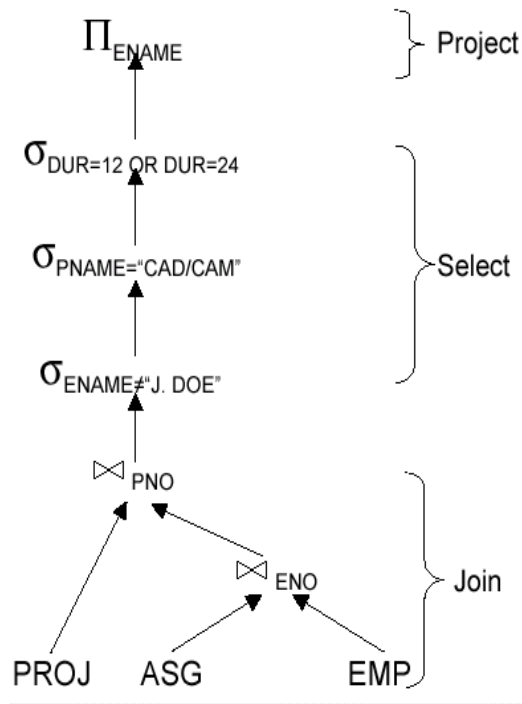
# Restructure

---

- To restructure the tree should be in a systematic way so that the “bad” operator trees are eliminated.
  - Allow the separation of the unary operations, simplifying the query expression
  - Unary operations on the same relations may be grouped so that access to a relation for performing unary operations can be done only once
  - Unary operations can be commuted with binary operations so that some operations may be done first
  - The binary operations can be ordered

Note: Generally, after the layer, the query is still far from providing an optimal execution, since information about data distribution and local fragments is not used at this layer

# Example of restructuring





---

## 6.2 Localization of Distributed Data



# Localization of Distributed Data

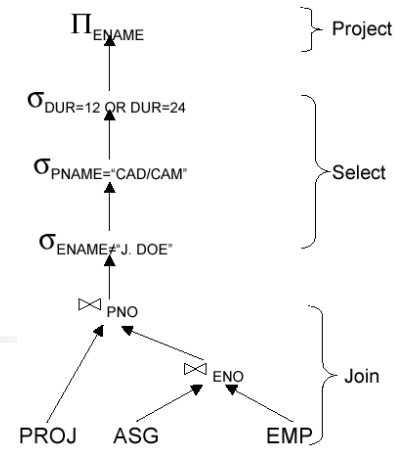
---

Input: Algebraic query on **global relations**

- Localize the query's data using data distributed information
  - Determine which fragments are involved
  - Transform the distributed (globe) query into a fragment query
- Step1: Localization program
  - substitute for each global query its materialization (reconstruction) program
- Step2: Optimize
  - simplification and restructuring

Output: Algebra query on **physical fragments**


# Example



Assume


- EMP is fragmented into  $EMP_1$ ,  $EMP_2$ ,  $EMP_3$  as follows:
  - ◆  $EMP_1 = \sigma_{ENO \leq "E3"}(EMP)$
  - ◆  $EMP_2 = \sigma_{"E3" < ENO \leq "E6"}(EMP)$
  - ◆  $EMP_3 = \sigma_{ENO \geq "E6"}(EMP)$
- ASG fragmented into  $ASG_1$  and  $ASG_2$  as follows:
  - ◆  $ASG_1 = \sigma_{ENO \leq "E3"}(ASG)$
  - ◆  $ASG_2 = \sigma_{ENO > "E3"}(ASG)$

Replace EMP by  $(EMP_1 \cup EMP_2 \cup EMP_3)$  and ASG by  $(ASG_1 \cup ASG_2)$  in any query






# Provides Parallelism



# Eliminates Unnecessary Work





## 6.2.1 Reduction for PHF

---

- The **localization program** for an PHF relation is the **union** of the fragments
- Reduction for PHF
  - After restructuring the subtrees, Determining those that will produce **empty relations**, and removing them

# Reduction with selection

- Selections on fragments that have a qualification contradicting the qualification of the fragmentation rule generate empty relations



Relation  $R$  and  $F_R = \{R_1, R_2, \dots, R_w\}$  where  $R_j = \sigma_{p_j}(R)$

$$\sigma_{p_i}(R_j) = \emptyset \text{ if } \forall x \text{ in } R: \neg(p_i(x) \wedge p_j(x))$$

Example

```
SELECT *  
FROM EMP  
WHERE ENO = "E5"
```

- $EMP_1 = \sigma_{ENO \leq "E3"}(EMP)$
- $EMP_2 = \sigma_{"E3" < ENO \leq "E6"}(EMP)$
- $EMP_3 = \sigma_{ENO \geq "E6"}(EMP)$





# Reduction with join

---

- Possible if fragmentation is done on join attribute
  - The simplification consists of distributing join over unions and eliminating useless joins

⇒ Distribute join over union

$$(R_1 \cup R_2) \bowtie S \Leftrightarrow (R_1 \bowtie S) \cup (R_2 \bowtie S)$$

⇒ Given  $R_i = \sigma_{p_i}(R)$  and  $S_j = \sigma_{p_j}(S)$

$$R_i \bowtie S_j = \emptyset \text{ if } \forall x \text{ in } R_i, \forall y \text{ in } S_j: \neg(p_i(x) \wedge p_j(y))$$

# Example of Reduction with join

Assume EMP is fragmented as before and

$ASG_1: \sigma_{ENO \leq "E3"}(ASG)$

$ASG_2: \sigma_{ENO > "E3"}(ASG)$

Consider the query

**SELECT \***


**FROM EMP, ASG**

**WHERE EMP.ENO=ASG.ENO**

♦  $EMP_1 = \sigma_{ENO \leq "E3"}(EMP)$


♦  $EMP_2 = \sigma_{"E3" < ENO \leq "E6"}(EMP)$

♦  $EMP_3 = \sigma_{ENO \geq "E6"}(EMP)$



# Example of Reduction with join

- Distribute join over unions
- Apply the reduction rule



- One advantage of the reduced query is that the partial joins can be done in parallel, and thus improve response time
- Reduced query is better when the number of partial joins is small



## 6.2.2 Reduction for VF

---

- The **localized program** for a VF relation consists of the **join** of the fragments on the common attribute
- Reduction for VF
  - Determining the **useless** intermediate relations and removing the subtrees that produce them
  - Projects on a vertical fragment that has no attributes in common with the projection attributes (except the key of the relation) produce useless, though not empty relations




# Example of reduction for VF

Relation  $R$  defined over attributes  $A = \{A_1, \dots, A_n\}$  vertically fragmented as  $R_i = \Pi_{A'}(R)$  where  $A' \subseteq A$ :

$\Pi_{D,K}(R_i)$  is useless if the set of projection attributes  $D$  is not in  $A'$

Example:  $EMP_1 = \Pi_{ENO,ENAME}(EMP)$ ;  $EMP_2 = \Pi_{ENO,TITLE}(EMP)$

```
SELECT  ENAME  
FROM    EMP
```





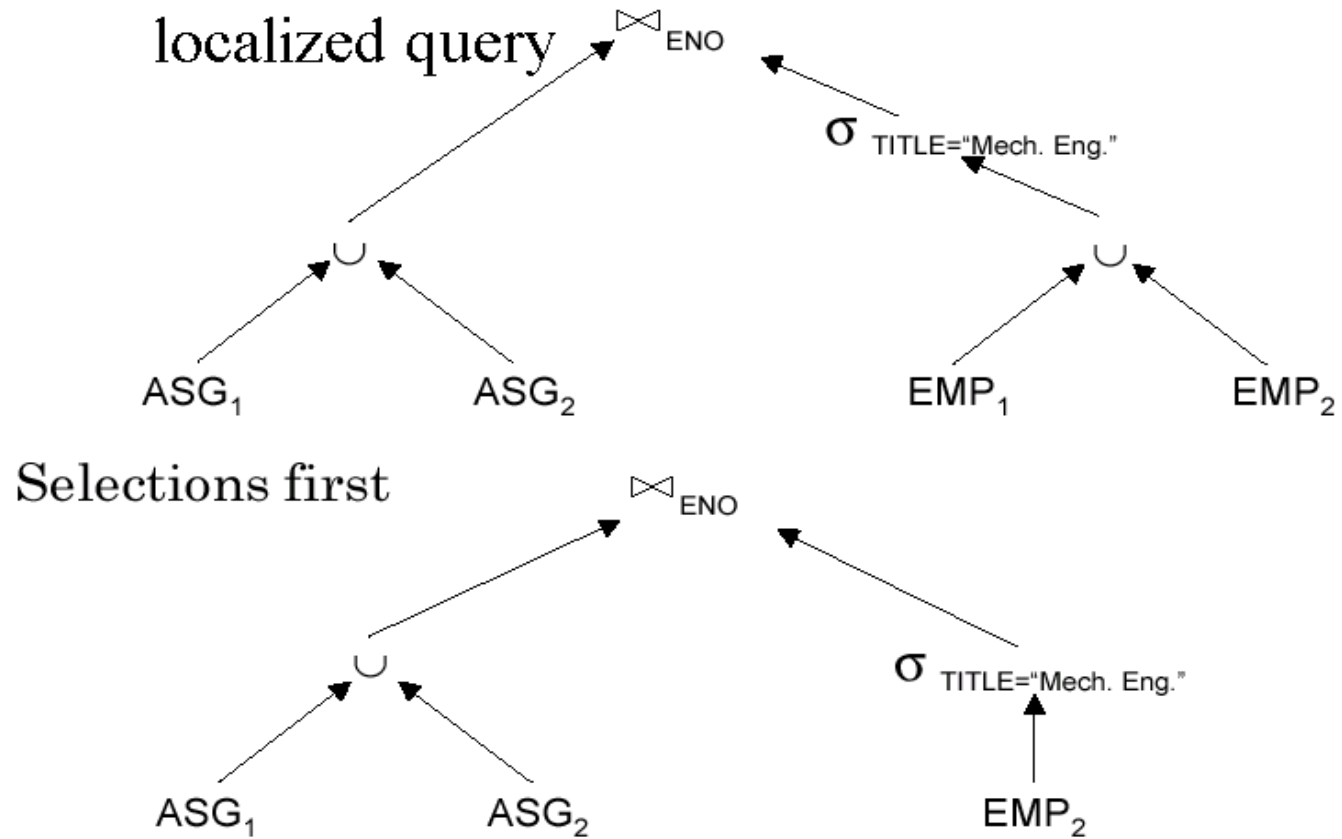
## 6.2.3 Reduction for DHF

---

- Rule :
  - Distribute joins over unions
  - Apply the join reduction for horizontal fragmentation
- Example
$$\text{ASG1 : ASG} \propto_{\text{ENO}} \text{EMP1}$$
$$\text{ASG2 : ASG} \propto_{\text{ENO}} \text{EMP2}$$
$$\text{EMP1 : } \delta_{\text{TITLE}=\text{"Programmer"}} (\text{EMP})$$
$$\text{EMP2 : } \delta_{\text{TITLE} \neq \text{"Programmer"}} (\text{EMP})$$
- Query


```
SELECT *
FROM EMP, ASG
WHERE ASG.ENO = EMP.ENO
AND EMP.TITLE = "Mech. Eng."
```

# Example of reduction for DHF




# Example of reduction for DHF

Joins over unions



Elimination of the empty intermediate relations  
(left sub-tree)



The reduced query is always preferable to the localized query because the number of partial joins usually equals the number of fragments



## 6.2.4 Reduction for HF

---

- Combine the rules already specified:
  - Remove empty relations generated by contradicting ***selections*** on horizontal fragments;
  - Remove useless relations generated by ***projections*** on vertical fragments;
  - Distribute joins over unions in order to isolate and remove useless ***joins***.

# Example of reduction for HF

Consider the following hybrid fragmentation:


$$EMP_1 = \sigma_{ENO \leq "E4"} (\Pi_{ENO, ENAME} (EMP))$$

$$EMP_2 = \sigma_{ENO > "E4"} (\Pi_{ENO, ENAME} (EMP))$$

$$EMP_3 = \Pi_{ENO, TITLE} (EMP)$$

and the query

```
SELECT      ENAME
FROM        EMP
WHERE       ENO = "E5"
```



Note: Generally, after the layer, the query is still far from providing an optimal execution



---

# Conclusion



# Conclusion

---

- Naive way
  - Query decomposition
  - Data localization
- Optimization in the above layers
  - Avoid the worse executions
- Global and local Optimization layers





# References

---

- M. Jarke and J. Koch. “Query Optimization in Database Systems,” *Computing Surveys*, June 1984, 16(2): 227-269.
- J. D. Ullman, Principles of Database Systems (2nd edition), Rockville, Md.: Computer Science Press, 1982.
- S. Ceri and Pelagatti, Correctness of Query Execution Strategies in Distributed Databases, ACM Trans. Database System (Dec. 1983), 8(4): 577-607